

Live CaseStudies 2025-11-12

November 12, 2025

1 Live Case Study 1 2025-11-12

Matriculation: 0815


```
[2]: import pandas as pd  
  
X = pd.read_csv("../PMCM_WS2425/Data02.csv")
```

```
[6]: X.describe()
```

```
[6]:          f1          f2      error  
count  10000.000000  10000.000000  10000.000000  
mean    -1.477134    2.914468    2.205596  
std     1.079116    1.019233    2.832223  
min    -5.973536   -1.874929    0.000158  
25%   -2.199008    2.260771    0.587542  
50%   -1.516600    2.912011    1.399142  
75%   -0.808268    3.580329    2.859412  
max     4.359184    6.409433   44.133879
```

1.1 Q5a

```
[4]: import matplotlib.pyplot as plt  
  
[8]: plt.hist(X.error,bins=50)  
plt.grid()
```



1.2 Q5b

```
[9]: X[X['error']>6].count()
```

```
[9]: f1      640  
f2      640  
error   640  
dtype: int64
```

640 bad parts

```
[10]: X[X['error']<=6].count()
```

```
[10]: f1      9360  
f2      9360  
error   9360  
dtype: int64
```

9360 good parts

```
[11]: #Production yield = good / (good + bad)  
9360 / (9360 + 640)
```

[11]: 0.936

1.3 Q5c

```
[14]: # we have a classification problem
# most stupid predictor would predict all as good (majority class)
# accuracy = correct / total = tp + tn / (tp + tn + fp + fn)

acc_stupid = 9360 / (9360 + 640)
print(f"The baseline is an accuracy of {acc_stupid}")
```

The baseline is an accuracy of 0.936

```
[20]: ## alternative way
from sklearn.dummy import DummyClassifier
from sklearn.metrics import accuracy_score

clf = DummyClassifier(strategy="most_frequent")

y_true = X['error']<=6
xx = X[['f1', 'f2']]

clf.fit(xx,y_true)

y_pred = clf.predict(xx)

acc = accuracy_score(y_true,y_pred)
print(f"Baseline accuracy {acc}")
```

Baseline accuracy 0.936

1.4 Q5d

The dataset is unbalanced, because the good class is represented in 93.6 % of the instances and the bad class only in 6.4 %.

Firstly we have to split in train and test data. Then train on the train and test on the test :-)

Because of imbalance we have to make sure that we get bad instances in train and in test!

One way to do that would be a stratified split (not just a normal split). Or we could * augment data with more bad cases (but how?? :-() * reduce number of good cases

1.5 Q6

```
[22]: X = pd.read_csv("../PMCM_WS2425/Data03.csv")
```

```
[25]: X.head()
```

```
[25]:      timestamp      f1      f2
0  2023-07-20 06:00:00 -5.143580 -5.382025
1  2023-07-20 06:01:00 -4.605052 -5.503890
2  2023-07-20 06:02:00 -5.690262 -5.285184
3  2023-07-20 06:03:00 -3.917527 -6.951911
4  2023-07-20 06:04:00 -4.788862 -5.260651
```

```
[26]: X['timestamp_parsed'] = pd.to_datetime(X.timestamp)
```

```
[27]: X.head()
```

```
[27]:      timestamp      f1      f2      timestamp_parsed
0  2023-07-20 06:00:00 -5.143580 -5.382025 2023-07-20 06:00:00
1  2023-07-20 06:01:00 -4.605052 -5.503890 2023-07-20 06:01:00
2  2023-07-20 06:02:00 -5.690262 -5.285184 2023-07-20 06:02:00
3  2023-07-20 06:03:00 -3.917527 -6.951911 2023-07-20 06:03:00
4  2023-07-20 06:04:00 -4.788862 -5.260651 2023-07-20 06:04:00
```

```
[28]: X.describe()
```

```
[28]:      f1      f2      timestamp_parsed
count  3840.000000  3840.000000          3840
mean    -4.861466   -4.854412  2023-07-22 01:59:30
min     -8.740101   -8.392300  2023-07-20 06:00:00
25%     -5.704431   -5.649969  2023-07-21 03:59:45
50%     -4.999363   -5.013370  2023-07-22 01:59:30
75%     -4.317683   -4.317863  2023-07-22 23:59:15
max      7.865602    7.463494  2023-07-23 21:59:00
std      1.604436    1.564089           NaN
```


```
[29]: print(f"Timestamp ranges from {X.timestamp_parsed.min()} to {X.timestamp_parsed.max()}" )
```

Timestamp ranges from 2023-07-20 06:00:00 to 2023-07-23 21:59:00

```
[30]: deltas = X.timestamp_parsed.diff()/pd.Timedelta(hours=1)
```

```
[31]: plt.plot(deltas)
```

```
[31]: [<matplotlib.lines.Line2D at 0x7f4193e5b6d0>]
```



```
[37]: X[deltas > 0.1].timestamp_parsed
```


```
[37]: 960    2023-07-21 06:00:00
1920    2023-07-22 06:00:00
2880    2023-07-23 06:00:00
Name: timestamp_parsed, dtype: datetime64[ns]
```

Yes, we have three breaks of 8 h each. We found them by subtracting `timestamp_i+1 - timestamp_i`

1.6 Q6b

```
[38]: plt.scatter(X.f1,X.f2)
```


```
[38]: <matplotlib.collections.PathCollection at 0x7f419370bcd0>
```



```
[41]: from sklearn.cluster import KMeans  
  
xx = X[['f1','f2']]  
  
clusterer = KMeans(n_clusters=2)  
  
clusterer.fit(xx)  
  
y_cluster = pd.Series(clusterer.predict(xx))
```

```
[45]: plt.scatter(xx.f1,xx.f2,c=y_cluster)
```

```
[45]: <matplotlib.collections.PathCollection at 0x7f4193e48880>
```



```
[46]: y_cluster[y_cluster == 0].count()
```

```
[46]: np.int64(3780)
```

```
[47]: y_cluster[y_cluster == 1].count()
```

```
[47]: np.int64(60)
```

```
[48]: X[y_cluster == 1].timestamp_parsed
```

```
[48]: 193    2023-07-20 09:13:00  
198    2023-07-20 09:18:00  
246    2023-07-20 10:06:00  
327    2023-07-20 11:27:00  
332    2023-07-20 11:32:00  
415    2023-07-20 12:55:00  
646    2023-07-20 16:46:00  
794    2023-07-20 19:14:00  
807    2023-07-20 19:27:00  
821    2023-07-20 19:41:00  
928    2023-07-20 21:28:00  
962    2023-07-21 06:02:00
```

968	2023-07-21 06:08:00
981	2023-07-21 06:21:00
1100	2023-07-21 08:20:00
1144	2023-07-21 09:04:00
1239	2023-07-21 10:39:00
1338	2023-07-21 12:18:00
1379	2023-07-21 12:59:00
1385	2023-07-21 13:05:00
1470	2023-07-21 14:30:00
1518	2023-07-21 15:18:00
1572	2023-07-21 16:12:00
1716	2023-07-21 18:36:00
1741	2023-07-21 19:01:00
1872	2023-07-21 21:12:00
1876	2023-07-21 21:16:00
1913	2023-07-21 21:53:00
2015	2023-07-22 07:35:00
2019	2023-07-22 07:39:00
2087	2023-07-22 08:47:00
2320	2023-07-22 12:40:00
2379	2023-07-22 13:39:00
2382	2023-07-22 13:42:00
2432	2023-07-22 14:32:00
2523	2023-07-22 16:03:00
2644	2023-07-22 18:04:00
2715	2023-07-22 19:15:00
2716	2023-07-22 19:16:00
2863	2023-07-22 21:43:00
2911	2023-07-23 06:31:00
2975	2023-07-23 07:35:00
3002	2023-07-23 08:02:00
3045	2023-07-23 08:45:00
3060	2023-07-23 09:00:00
3207	2023-07-23 11:27:00
3224	2023-07-23 11:44:00
3250	2023-07-23 12:10:00
3287	2023-07-23 12:47:00
3431	2023-07-23 15:11:00
3509	2023-07-23 16:29:00
3539	2023-07-23 16:59:00
3553	2023-07-23 17:13:00
3559	2023-07-23 17:19:00
3622	2023-07-23 18:22:00
3625	2023-07-23 18:25:00
3728	2023-07-23 20:08:00
3759	2023-07-23 20:39:00
3761	2023-07-23 20:41:00

3794 2023-07-23 21:14:00

Name: timestamp_parsed, dtype: datetime64[ns]

[]: